

# **Distributed Systems**

(3rd Edition)

## Chapter 02: Architectures

Version: August 29, 2017

# Architectural styles

## Basic idea

A style is formulated in terms of

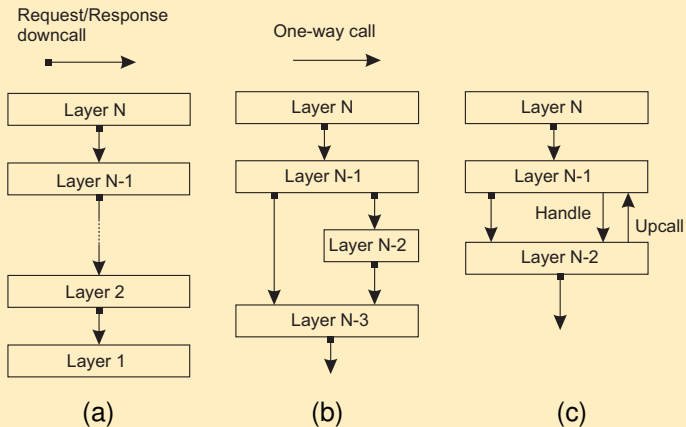
- (replaceable) components with well-defined interfaces
- the way that components are connected to each other
- the data exchanged between components
- how these components and connectors are jointly configured into a system.

## Connector

A mechanism that mediates communication, coordination, or cooperation among components. **Example:** facilities for (remote) procedure call, messaging, or streaming.


# Layered architecture

## Different layered organizations



# Example: communication protocols

## Protocol, service, interface



# Two-party communication

## Server

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 (conn, addr) = s.accept() # returns new socket and addr. client
4 while True:               # forever
5     data = conn.recv(1024) # receive data from client
6     if not data: break     # stop if client stopped
7     conn.send(str(data)+"*") # return sent data plus an "*"
8     conn.close()          # close the connection
```

## Client

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 s.connect((HOST, PORT)) # connect to server (block until accepted)
4 s.send('Hello, world') # send some data
5 data = s.recv(1024)    # receive the response
6 print data             # print the result
7 s.close()              # close the connection
```

# Application Layering

## Traditional three-layered view

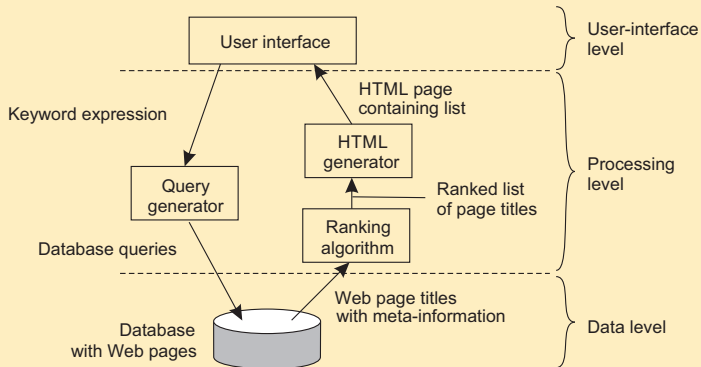
- **Application-interface layer** contains units for interfacing to users or external applications
- **Processing layer** contains the functions of an application, i.e., without specific data
- **Data layer** contains the data that a client wants to manipulate through the application components

## Observation

This layering is found in many distributed information systems, using traditional database technology and accompanying applications.

# Application Layering



## Example: a simple search engine



# Object-based style

## Essence

Components are objects, connected to each other through procedure calls. Objects may be placed on different machines; calls can thus execute across a network.



## Encapsulation

Objects are said to **encapsulate data** and offer **methods on that data** without revealing the internal implementation.



# RESTful architectures

## Essence

View a distributed system as a collection of resources, individually managed by components. Resources may be added, removed, retrieved, and modified by (remote) applications.

- 1 Resources are identified through a single naming scheme
- 2 All services offer the same interface
- 3 Messages sent to or from a service are fully self-described
- 4 After executing an operation at a service, that component forgets everything about the caller

## Basic operations

Operation	Description
PUT	Create a new resource
GET	Retrieve the state of a resource in some representation
DELETE	Delete a resource
POST	Modify a resource by transferring a new state

# Example: Amazon's Simple Storage Service

## Essence

**Objects** (i.e., files) are placed into **buckets** (i.e., directories). Buckets cannot be placed into buckets. Operations on `ObjectName` in bucket `BucketName` require the following identifier:

```
http://BucketName.s3.amazonaws.com/ObjectName
```

## Typical operations

All operations are carried out by sending HTTP requests:

- Create a bucket/object: `PUT`, along with the URI
- Listing objects: `GET` on a bucket name
- Reading an object: `GET` on a full URI

# On interfaces

## Issue

Many people like RESTful approaches because the interface to a service is so simple. The catch is that much needs to be done in the **parameter space**.

## Amazon S3 SOAP interface

Bucket operations	Object operations
ListAllMyBuckets	PutObjectInline
CreateBucket	PutObject
DeleteBucket	CopyObject
ListBucket	GetObject
GetBucketAccessControlPolicy	GetObjectExtended
SetBucketAccessControlPolicy	DeleteObject
GetBucketLoggingStatus	GetObjectAccessControlPolicy
SetBucketLoggingStatus	SetObjectAccessControlPolicy

# On interfaces

## Simplifications

Assume an interface `bucket` offering an operation `create`, requiring an input string such as `mybucket`, for creating a bucket “mybucket.”

## SOAP

```
import bucket  
bucket.create("mybucket")
```

## RESTful

```
PUT "http://mybucket.s3.amazonaws.com/"
```

## Conclusions



Are there any to draw?

# Coordination

## Temporal and referential coupling

	Temporally coupled	Temporally decoupled
Referentially coupled	Direct	Mailbox
Referentially decoupled	Event-based	Shared data space

## Event-based and Shared data space



# Example: Linda tuple space

## Three simple operations

- $\text{in}(t)$ : remove a tuple matching template  $t$
- $\text{rd}(t)$ : obtain copy of a tuple matching template  $t$
- $\text{out}(t)$ : add tuple  $t$  to the tuple space

## More details

- Calling  $\text{out}(t)$  twice in a row, leads to storing **two** copies of tuple  $t \Rightarrow$  a tuple space is modeled as a **multiset**.
- Both  $\text{in}$  and  $\text{rd}$  are **blocking** operations: the caller will be blocked until a matching tuple is found, or has become available.

# Example: Linda tuple space

## Bob

```
1 blog = linda.universe._rd(("MicroBlog",linda.TupleSpace))[1]
2
3 blog._out(("bob","distsys","I am studying chap 2"))
4 blog._out(("bob","distsys","The linda example's pretty simple"))
5 blog._out(("bob","gtcn","Cool book!"))
```

## Alice

```
1 blog = linda.universe._rd(("MicroBlog",linda.TupleSpace))[1]
2
3 blog._out(("alice","gtcn","This graph theory stuff is not easy"))
4 blog._out(("alice","distsys","I like systems more than graphs"))
```

## Chuck

```
1 blog = linda.universe._rd(("MicroBlog",linda.TupleSpace))[1]
2
3 t1 = blog._rd(("bob","distsys",str))
4 t2 = blog._rd(("alice","gtcn",str))
5 t3 = blog._rd(("bob","gtcn",str))
```

# Using legacy to build middleware

## Problem

The interfaces offered by a legacy component are most likely not suitable for all applications.



## Solution

A **wrapper** or **adapter** offers an interface acceptable to a client application. Its functions are transformed into those available at the component.



# Organizing wrappers

Two solutions: 1-on-1 or through a broker



Complexity with  $N$  applications

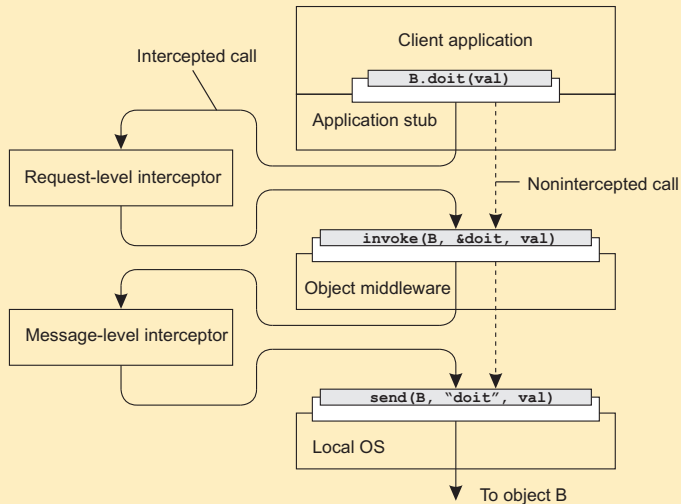
- **1-on-1**: requires  $N \times (N - 1) = \mathcal{O}(N^2)$  wrappers
- **broker**: requires  $2N = \mathcal{O}(N)$  wrappers

# Developing adaptable middleware

## Problem

Middleware contains solutions that are good for **most** applications  $\Rightarrow$  you may want to adapt its behavior for specific applications.

# Intercept the usual flow of control




# Centralized system architectures

## Basic Client–Server Model

### Characteristics:

- There are processes offering services (**servers**)
- There are processes that use services (**clients**)
- Clients and servers can be on different machines
- Clients follow request/reply model with respect to using services

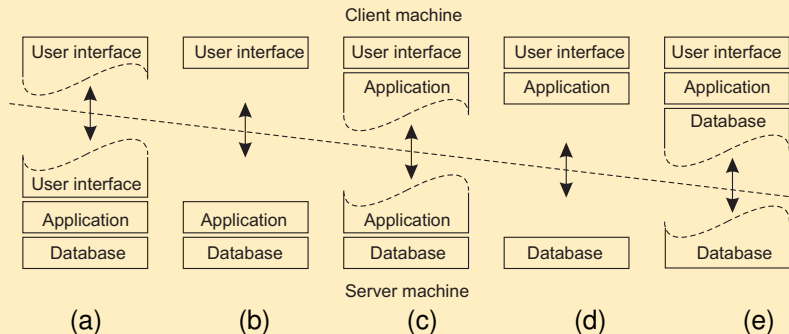


# Multi-tiered centralized system architectures

## Some traditional organizations


- **Single-tiered:** dumb terminal/mainframe configuration
- **Two-tiered:** client/single server configuration
- **Three-tiered:** each layer on separate machine

## Traditional two-tiered configurations



# Being client and server at the same time

## Three-tiered architecture



# Alternative organizations

## Vertical distribution

Comes from dividing distributed applications into three logical layers, and running the components from each layer on a different server (machine).

## Horizontal distribution

A client or server may be physically split up into logically equivalent parts, but each part is operating on its own share of the complete data set.

## Peer-to-peer architectures

Processes are all equal: the functions that need to be carried out are represented by every process  $\Rightarrow$  each process will act as a client and a server at the same time (i.e., acting as a **servant**).

# Structured P2P


## Essence

Make use of a **semantic-free index**: each data item is uniquely associated with a key, in turn used as an index. Common practice: use a **hash function**

$$\text{key}(\text{data item}) = \text{hash}(\text{data item's value}).$$

P2P system now responsible for storing  $(\text{key}, \text{value})$  pairs.

## Simple example: hypercube



Looking up  $d$  with **key**  $k \in \{0, 1, 2, \dots, 2^4 - 1\}$  means **routing** request to node with **identifier**  $k$ .




# Example: Chord

## Principle

- Nodes are logically organized in a ring. Each node has an  $m$ -bit **identifier**.
- Each data item is hashed to an  $m$ -bit **key**.
- Data item with key  $k$  is stored at node with smallest identifier  $id \geq k$ , called the **successor** of key  $k$ .
- The ring is extended with various **shortcut links** to other nodes.

# Example: Chord



# Unstructured P2P

## Essence

Each node maintains an ad hoc list of neighbors. The resulting overlay resembles a **random graph**: an edge  $\langle u, v \rangle$  exists only with a certain probability  $\mathbb{P}[\langle u, v \rangle]$ .

## Searching

- **Flooding**: issuing node  $u$  passes request for  $d$  to all neighbors. Request is ignored when receiving node had seen it before. Otherwise,  $v$  searches locally for  $d$  (recursively). May be limited by a **Time-To-Live**: a maximum number of hops.
- **Random walk**: issuing node  $u$  passes request for  $d$  to randomly chosen neighbor,  $v$ . If  $v$  does not have  $d$ , it forwards request to one of *its* randomly chosen neighbors, and so on.

# Flooding versus random walk

## Model

Assume  $N$  nodes and that each data item is replicated across  $r$  randomly chosen nodes.

## Random walk

$\mathbb{P}[k]$  probability that item is found after  $k$  attempts:

$$\mathbb{P}[k] = \frac{r}{N} \left(1 - \frac{r}{N}\right)^{k-1}.$$

$S$  (“search size”) is expected number of nodes that need to be probed:

$$S = \sum_{k=1}^N k \cdot \mathbb{P}[k] = \sum_{k=1}^N k \cdot \frac{r}{N} \left(1 - \frac{r}{N}\right)^{k-1} \approx N/r \text{ for } 1 \ll r \leq N.$$

# Flooding versus random walk

## Flooding

- Flood to  $d$  randomly chosen neighbors
- After  $k$  steps, some  $R(k) = d \cdot (d-1)^{k-1}$  will have been reached (assuming  $k$  is small).
- With fraction  $r/N$  nodes having data, if  $\frac{r}{N} \cdot R(k) \geq 1$ , we will have found the data item.

## Comparison


- If  $r/N = 0.001$ , then  $S \approx 1000$
- With flooding and  $d = 10, k = 4$ , we contact 7290 nodes.
- Random walks are more communication efficient, but might take longer before they find the result.

# Super-peer networks

## Essence

It is sometimes sensible to break the symmetry in pure peer-to-peer networks:

- When searching in unstructured P2P systems, having **index servers** improves performance
- Deciding where to store data can often be done more efficiently through **brokers**.



# Skype's principle operation: *A* wants to contact *B*

Both *A* and *B* are on the public Internet

- A TCP connection is set up between *A* and *B* for control packets.
- The actual call takes place using UDP packets between negotiated ports.

*A* operates behind a firewall, while *B* is on the public Internet

- *A* sets up a TCP connection (for control packets) to a super peer *S*
- *S* sets up a TCP connection (for relaying control packets) to *B*
- The actual call takes place through UDP and directly between *A* and *B*


Both *A* and *B* operate behind a firewall

- *A* connects to an online super peer *S* through TCP
- *S* sets up TCP connection to *B*.
- For the actual call, another super peer is contacted to act as a **relay** *R*: *A* sets up a connection to *R*, and so will *B*.
- All voice traffic is forwarded over the two TCP connections, and through *R*.

# Edge-server architecture

## Essence

Systems deployed on the Internet where servers are placed **at the edge** of the network: the boundary between enterprise networks and the actual Internet.






# Collaboration: The BitTorrent case

## Principle: search for a file $F$

- Lookup file at a global directory  $\Rightarrow$  returns a **torrent file**
- Torrent file contains reference to **tracker**: a server keeping an accurate account of **active** nodes that have (chunks of)  $F$ .
- $P$  can join **swarm**, get a chunk for free, and then trade a copy of that chunk for another one with a peer  $Q$  also in the swarm.



# BitTorrent under the hood

## Some essential details

- A tracker for file  $F$  returns the set of its downloading processes: the current **swarm**.
- $A$  communicates only with a subset of the swarm: the **neighbor set**  $N_A$ .
- if  $B \in N_A$  then also  $A \in N_B$ .
- Neighbor sets are regularly updated by the tracker

## Exchange blocks

- A file is divided into equally sized **pieces** (typically each being 256 KB)
- Peers exchange **blocks** of pieces, typically some 16 KB.
- $A$  can upload a block  $d$  of piece  $D$ , only if it has piece  $D$ .
- Neighbor  $B$  belongs to the **potential set**  $P_A$  of  $A$ , if  $B$  has a block that  $A$  needs.
- If  $B \in P_A$  and  $A \in P_B$ :  $A$  and  $B$  are in a position that they can **trade** a block.

# BitTorrent phases

## Bootstrap phase

A has just received its first piece (through **optimistic unchoking**: a node from  $N_A$  unselfishly provides the blocks of a piece to get a newly arrived node started).

## Trading phase

$|P_A| > 0$ : there is (in principle) always a peer with whom A can trade.

## Last download phase

$|P_A| = 0$ : A is dependent on newly arriving peers in  $N_A$  in order to get the last missing pieces.  $N_A$  can change only through the tracker.

# BitTorrent phases

Development of  $|P|$  relative to  $|N|$ .

