

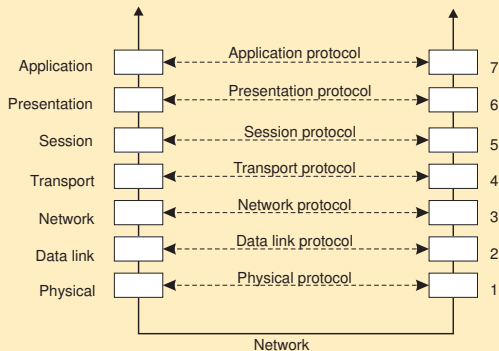
# **Distributed Systems**

(3rd Edition)

## Chapter 04: Communication

Version: August 29, 2017

# Basic networking model



## Drawbacks

- Focus on message-passing only
- Often unneeded or unwanted functionality
- Violates access transparency

# Low-level layers

## Recap

- **Physical layer**: contains the specification and implementation of bits, and their transmission between sender and receiver
- **Data link layer**: prescribes the transmission of a series of bits into a frame to allow for error and flow control
- **Network layer**: describes how packets in a network of computers are to be **routed**.

## Observation

For many distributed systems, the lowest-level interface is that of the network layer.

# Transport Layer

## Important

The transport layer provides the actual communication facilities for most distributed systems.

## Standard Internet protocols

- TCP: connection-oriented, reliable, stream-oriented communication
- UDP: unreliable (best-effort) datagram communication

# Middleware layer

## Observation


Middleware is invented to provide **common** services and protocols that can be used by many **different** applications

- A rich set of **communication protocols**
- **(Un)marshaling** of data, necessary for integrated systems
- **Naming protocols**, to allow easy sharing of resources
- **Security protocols** for secure communication
- **Scaling mechanisms**, such as for replication and caching

## Note


What remains are truly **application-specific** protocols... **such as?**

# An adapted layering scheme



# Types of communication


## Distinguish...



- Transient versus persistent communication
- Asynchronous versus synchronous communication

# Types of communication

## Transient versus persistent




- **Transient communication:** Comm. server discards message when it cannot be delivered at the next server, or at the receiver.
- **Persistent communication:** A message is stored at a communication server as long as it takes to deliver it.



# Types of communication

## Places for synchronization



- At request submission
- At request delivery
- After request processing

# Client/Server

## Some observations

Client/Server computing is generally based on a model of **transient synchronous communication**:

- Client and server have to be active at time of communication
- Client issues request and blocks until it receives reply
- Server essentially waits only for incoming requests, and subsequently processes them

## Drawbacks synchronous communication

- Client cannot do any other work while waiting for reply
- Failures have to be handled immediately: the client is waiting
- The model may simply not be appropriate (mail, news)

# Messaging

## Message-oriented middleware

Aims at high-level **persistent asynchronous communication**:

- Processes send each other messages, which are queued
- Sender need not wait for immediate reply, but can do other things
- Middleware often ensures fault tolerance


# Basic RPC operation

## Observations

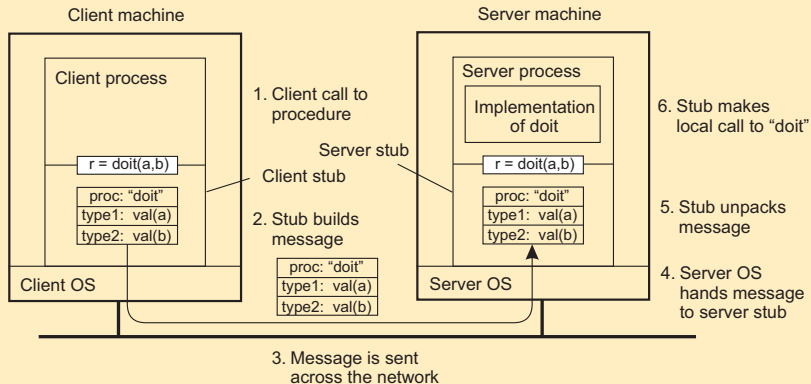
- Application developers are familiar with simple procedure model
- Well-engineered procedures operate in isolation (black box)
- There is no fundamental reason not to execute procedures on separate machine

## Conclusion

Communication between caller & callee can be hidden by using procedure-call mechanism.



# Basic RPC operation



- 1 Client procedure calls client stub.
- 2 Stub builds message; calls local OS.
- 3 OS sends message to remote OS.
- 4 Remote OS gives message to stub.
- 5 Stub unpacks parameters; calls server.
- 6 Server does local call; returns result to stub.
- 7 Stub builds message; calls OS.
- 8 OS sends message to client's OS.
- 9 Client's OS gives message to stub.
- 10 Client stub unpacks result; returns to client.

# RPC: Parameter passing

There's more than just wrapping parameters into a message

- Client and server machines may have **different data representations** (think of byte ordering)
- Wrapping a parameter means **transforming a value into a sequence of bytes**
- Client and server have to **agree on the same encoding**:
  - How are **basic data values** represented (integers, floats, characters)
  - How are **complex data values** represented (arrays, unions)

## Conclusion

Client and server need to **properly interpret messages**, transforming them into machine-dependent representations.

# RPC: Parameter passing

## Some assumptions

- **Copy in/copy out** semantics: while procedure is executed, nothing can be assumed about parameter values.
- **All** data that is to be operated on is passed by parameters. Excludes passing **references to (global) data**.

## Conclusion

Full access transparency cannot be realized.


A **remote reference** mechanism enhances access transparency

- Remote reference offers **unified access** to remote data
- Remote references can be **passed as parameter** in RPCs
- **Note**: stubs can sometimes be used as such references

# Asynchronous RPCs

## Essence

Try to get rid of the strict request-reply behavior, but let the client continue without waiting for an answer from the server.

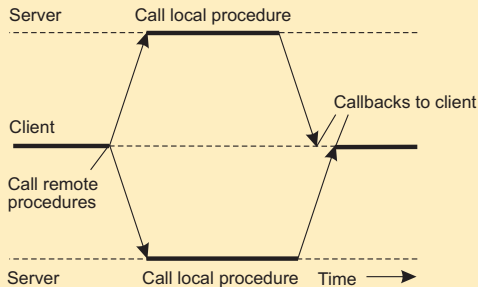




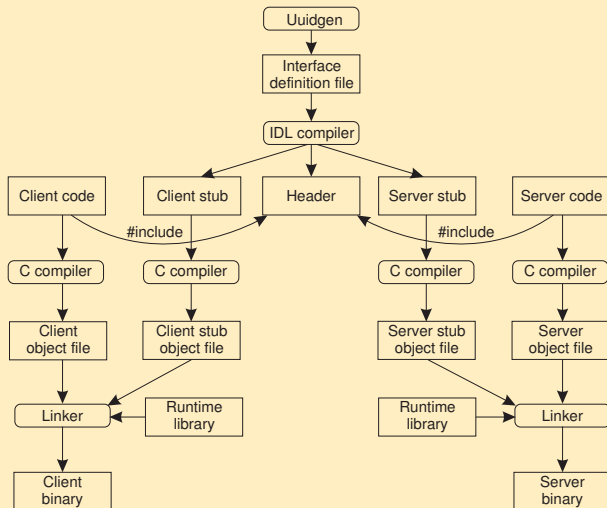
# Sending out multiple RPCs

## Essence

Sending an RPC request to a group of servers.



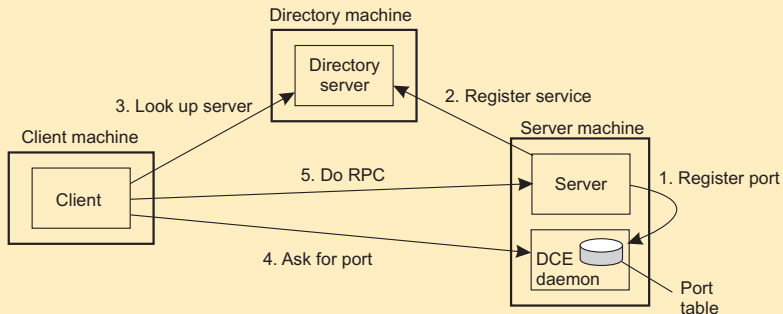
# RPC in practice



# Client-to-server binding (DCE)

## Issues


(1) Client must locate server machine, and (2) locate the server.



# Transient messaging: sockets

## Berkeley socket interface

Operation	Description
socket	Create a new communication end point
bind	Attach a local address to a socket
listen	Tell operating system what the maximum number of pending connection requests should be
accept	Block caller until a connection request arrives
connect	Actively attempt to establish a connection
send	Send some data over the connection
receive	Receive some data over the connection
close	Release the connection



# Sockets: Python code

## Server

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 s.bind((HOST, PORT))
4 s.listen(1)
5 (conn, addr) = s.accept() # returns new socket and addr. client
6 while True:               # forever
7     data = conn.recv(1024) # receive data from client
8     if not data: break     # stop if client stopped
9     conn.send(str(data)+"*") # return sent data plus an "*"
10 conn.close()              # close the connection
```

## Client

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 s.connect((HOST, PORT)) # connect to server (block until accepted)
4 s.send('Hello, world') # send same data
5 data = s.recv(1024)    # receive the response
6 print data             # print the result
7 s.close()               # close the connection
```

# Making sockets easier to work with

## Observation

Sockets are rather low level and programming mistakes are easily made. However, the way that they are used is often the same (such as in a client-server setting).

## Alternative: ZeroMQ

Provides a higher level of expression by **pairing** sockets: one for sending messages at process  $P$  and a corresponding one at process  $Q$  for receiving messages. All communication is **asynchronous**.

## Three patterns

- Request-reply
- Publish-subscribe
- Pipeline

# Request-reply

## Server

```
1 import zmq
2 context = zmq.Context()
3
4 p1 = "tcp://" + HOST + ":" + PORT1 # how and where to connect
5 p2 = "tcp://" + HOST + ":" + PORT2 # how and where to connect
6 s = context.socket(zmq.REP)        # create reply socket
7
8 s.bind(p1)                          # bind socket to address
9 s.bind(p2)                          # bind socket to address
10 while True:
11     message = s.recv()              # wait for incoming message
12     if not "STOP" in message:      # if not to stop...
13         s.send(message + "*")      # append "*" to message
14     else:                           # else...
15         break                       # break out of loop and end
```

# Request-reply

## Client

```
1 import zmq
2 context = zmq.Context()
3
4 php = "tcp://" + HOST + ":" + PORT # how and where to connect
5 s = context.socket(zmq.REQ) # create socket
6
7 s.connect(php) # block until connected
8 s.send("Hello World") # send message
9 message = s.recv() # block until response
10 s.send("STOP") # tell server to stop
11 print message # print result
```



# Publish-subscribe

## Server

```
1 import zmq, time
2
3 context = zmq.Context()
4 s = context.socket(zmq.PUB)           # create a publisher socket
5 p = "tcp://" + HOST + ":" + PORT     # how and where to communicate
6 s.bind(p)                             # bind socket to the address
7 while True:
8     time.sleep(5)                     # wait every 5 seconds
9     s.send("TIME " + time.asctime())  # publish the current time
```

## Client

```
1 import zmq
2
3 context = zmq.Context()
4 s = context.socket(zmq.SUB)           # create a subscriber socket
5 p = "tcp://" + HOST + ":" + PORT     # how and where to communicate
6 s.connect(p)                          # connect to the server
7 s.setsockopt(zmq.SUBSCRIBE, "TIME")  # subscribe to TIME messages
8
9 for i in range(5):                    # Five iterations
10     time = s.recv()                  # receive a message
11     print time
```

# Pipeline

## Source

```
1 import zmq, time, pickle, sys, random
2
3 context = zmq.Context()
4 me = str(sys.argv[1])
5 s = context.socket(zmq.PUSH)           # create a push socket
6 src = SRC1 if me == '1' else SRC2      # check task source host
7 prt = PORT1 if me == '1' else PORT2    # check task source port
8 p = "tcp://" + src + ":" + prt         # how and where to connect
9 s.bind(p)                             # bind socket to address
10
11 for i in range(100):                  # generate 100 workloads
12     workload = random.randint(1, 100) # compute workload
13     s.send(pickle.dumps((me, workload))) # send workload to worker
```

# Pipeline

## Worker

```
1 import zmq, time, pickle, sys
2
3 context = zmq.Context()
4 me = str(sys.argv[1])
5 r = context.socket(zmq.PULL)      # create a pull socket
6 p1 = "tcp://" + SRC1 + ":" + PORT1 # address first task source
7 p2 = "tcp://" + SRC2 + ":" + PORT2 # address second task source
8 r.connect(p1)                    # connect to task source 1
9 r.connect(p2)                    # connect to task source 2
10
11 while True:
12     work = pickle.loads(r.recv()) # receive work from a source
13     time.sleep(work[1]*0.01)      # pretend to work
```

# MPI: When lots of flexibility is needed

## Representative operations

Operation	Description
<code>MPI_bsend</code>	Append outgoing message to a local send buffer
<code>MPI_send</code>	Send a message and wait until copied to local or remote buffer
<code>MPI_ssend</code>	Send a message and wait until transmission starts
<code>MPI_sendrecv</code>	Send a message and wait for reply
<code>MPI_issend</code>	Pass reference to outgoing message, and continue
<code>MPI_issend</code>	Pass reference to outgoing message, and wait until receipt starts
<code>MPI_recv</code>	Receive a message; block if there is none
<code>MPI_irecv</code>	Check if there is an incoming message, but do not block

# Message-oriented middleware

## Essence

Asynchronous persistent communication through support of middleware-level **queues**. Queues correspond to buffers at communication servers.

## Operations


Operation	Description
put	Append a message to a specified queue
get	Block until the specified queue is nonempty, and remove the first message
poll	Check a specified queue for messages, and remove the first. Never block
notify	Install a handler to be called when a message is put into the specified queue

# General model

## Queue managers

Queues are managed by **queue managers**. An application can put messages only into a **local** queue. Getting a message is possible by extracting it from a **local** queue only  $\Rightarrow$  queue managers need to **route** messages.

## Routing



# Message broker


## Observation

Message queuing systems assume a **common messaging protocol**: all applications agree on message format (i.e., structure and data representation)

## Broker handles application heterogeneity in an MQ system

- Transforms incoming messages to target format
- Very often acts as an **application gateway**
- May provide **subject-based** routing capabilities (i.e., **publish-subscribe** capabilities)

# Message broker: general architecture





# IBM's WebSphere MQ

## Basic concepts

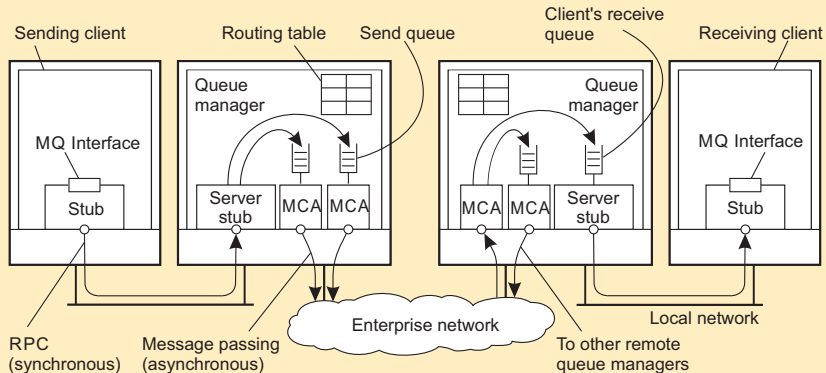
- **Application-specific messages** are put into, and removed from **queues**
- Queues reside under the regime of a **queue manager**
- Processes can put messages only in local queues, or through an RPC mechanism

## Message transfer

- Messages are transferred between queues
- Message transfer between queues at different processes, requires a **channel**
- At each end point of channel is a **message channel agent**
- Message channel agents are responsible for:
  - Setting up channels using lower-level network communication facilities (e.g., TCP/IP)
  - (Un)wrapping messages from/in transport-level packets
  - Sending/receiving packets

# IBM's WebSphere MQ

## Schematic overview



- Channels are inherently unidirectional
- Automatically start MCAs when messages arrive
- Any network of queue managers can be created
- Routes are set up manually (system administration)

# Message channel agents

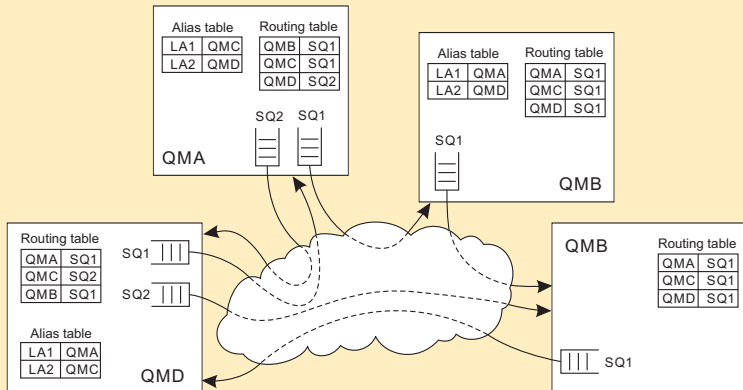
## Some attributes associated with message channel agents

Attribute	Description
Transport type	Determines the transport protocol to be used
FIFO delivery	Indicates that messages are to be delivered in the order they are sent
Message length	Maximum length of a single message
Setup retry count	Specifies maximum number of retries to start up the remote MCA
Delivery retries	Maximum times MCA will try to put received message into queue

# IBM's WebSphere MQ

## Routing

By using **logical names**, in combination with name resolution to local queues, it is possible to put a message in a **remote queue**



# Application-level multicasting

## Essence

Organize nodes of a distributed system into an **overlay network** and use that network to disseminate data:

- Oftentimes a **tree**, leading to unique paths
- Alternatively, also **mesh networks**, requiring a form of **routing**


# Application-level multicasting in Chord

## Basic approach

- ➊ Initiator generates a **multicast identifier**  $mid$ .
- ➋ Lookup  $succ(mid)$ , the node responsible for  $mid$ .
- ➌ Request is routed to  $succ(mid)$ , which will become the **root**.
- ➍ If  $P$  wants to join, it sends a **join** request to the root.
- ➎ When request arrives at  $Q$ :
  - $Q$  has not seen a join request before  $\Rightarrow$  it becomes **forwarder**;  $P$  becomes child of  $Q$ . **Join request continues to be forwarded.**
  - $Q$  knows about tree  $\Rightarrow P$  becomes child of  $Q$ . **No need to forward join request anymore.**

# ALM: Some costs

## Different metrics



- **Link stress:** How often does an ALM message cross the same physical link? **Example:** message from A to D needs to cross  $\langle Ra, Rb \rangle$  twice.
- **Stretch:** Ratio in delay between ALM-level path and network-level path. **Example:** messages B to C follow path of length 73 at ALM, but 47 at network level  $\Rightarrow$  stretch =  $73/47$ .

# Flooding


## Essence

$P$  simply sends a message  $m$  to each of its neighbors. Each neighbor will forward that message, except to  $P$ , and only if it had not seen  $m$  before.

## Performance

The more edges, the more expensive!

## The size of a random overlay as function of the number of nodes



## Variation

Let  $Q$  forward a message with a certain probability  $p_{\text{flood}}$ , possibly even dependent on its own number of neighbors (i.e., **node degree**) or the degree of its neighbors.



# Epidemic protocols

## Assume there are no write–write conflicts

- Update operations are performed at a single server
- A replica passes updated state to only a few neighbors
- Update propagation is lazy, i.e., not immediate
- Eventually, each update should reach every replica

## Two forms of epidemics

- **Anti-entropy**: Each replica regularly chooses another replica at random, and exchanges state differences, leading to identical states at both afterwards
- **Rumor spreading**: A replica which has just been updated (i.e., has been **contaminated**), tells a number of other replicas about its update (contaminating them as well).

# Anti-entropy

## Principle operations

- A node  $P$  selects another node  $Q$  from the system at random.
- **Pull**:  $P$  only pulls in new updates from  $Q$
- **Push**:  $P$  only pushes its own updates to  $Q$
- **Push-pull**:  $P$  and  $Q$  send updates to each other

## Observation

For push-pull it takes  $\mathcal{O}(\log(N))$  rounds to disseminate updates to all  $N$  nodes (**round** = when every node has taken the initiative to start an exchange).


# Anti-entropy: analysis

## Basics


Consider a single source, propagating its update. Let  $p_i$  be the probability that a node has not received the update after the  $i^{\text{th}}$  round.

## Analysis: staying ignorant

- With **pull**,  $p_{i+1} = (p_i)^2$ : the node was not updated during the  $i^{\text{th}}$  round and should contact another ignorant node during the next round.
- With **push**,  
 $p_{i+1} = p_i(1 - \frac{1}{N})^{N(1-p_i)} \approx p_i e^{-1}$  (for small  $p_i$  and large  $N$ ): the node was ignorant during the  $i^{\text{th}}$  round and no updated node chooses to contact it during the next round.
- With **push-pull**:  $(p_i)^2 \cdot (p_i e^{-1})$



# Anti-entropy performance



# Rumor spreading

## Basic model

A server  $S$  having an update to report, contacts other servers. If a server is contacted to which the update has already propagated,  $S$  stops contacting other servers with probability  $p_{stop}$ .

## Observation

If  $s$  is the fraction of ignorant servers (i.e., which are unaware of the update), it can be shown that with many servers

$$s = e^{-(1/p_{stop}+1)(1-s)}$$

# Formal analysis

## Notations

Let  $s$  denote fraction of nodes that have not yet been updated (i.e., **susceptible**;  $i$  the fraction of updated (**infected**) and active nodes; and  $r$  the fraction of updated nodes that gave up (**removed**).

## From theory of epidemics

$$(1) \quad ds/dt = -s \cdot i$$

$$(2) \quad di/dt = s \cdot i - p_{stop} \cdot (1 - s) \cdot i$$

$$\Rightarrow \quad di/ds = -(1 + p_{stop}) + \frac{p_{stop}}{s}$$


$$\Rightarrow \quad i(s) = -(1 + p_{stop}) \cdot s + p_{stop} \cdot \ln(s) + C$$

## Wrapup

$i(1) = 0 \Rightarrow C = 1 + p_{stop} \Rightarrow i(s) = (1 + p_{stop}) \cdot (1 - s) + p_{stop} \cdot \ln(s)$ . We are looking for the case  $i(s) = 0$ , which leads to  $s = e^{-(1/p_{stop}+1)(1-s)}$

# Rumor spreading

## The effect of stopping



Consider 10,000 nodes

$1/p_{stop}$	$s$	$N_s$
1	0.203188	2032
2	0.059520	595
3	0.019827	198
4	0.006977	70
5	0.002516	25
6	0.000918	9
7	0.000336	3

### Note

If we really have to ensure that all servers are eventually updated, rumor spreading alone is not enough

# Deleting values

## Fundamental problem

We cannot remove an old value from a server and expect the removal to propagate. Instead, mere removal will be undone in due time using epidemic algorithms

## Solution

Removal has to be registered as a special update by inserting a **death certificate**



# Deleting values

When to remove a death certificate (it is not allowed to stay for ever)

- Run a global algorithm to detect whether the removal is known everywhere, and then collect the death certificates (looks like garbage collection)
- Assume death certificates propagate in finite time, and associate a maximum lifetime for a certificate (can be done at risk of not reaching all servers)

## Note

It is necessary that a removal actually reaches all servers.